

# Neural networks

## Cours 1/3


**Machine Learning**  
ModIA 2022

*Enseignant: Sixin Zhang*  
`sixin.zhang@toulouse-inp.fr`

- *Pattern Recognition and Machine Learning*, Christopher M. Bishop - 2006
- *Deep Learning*, I Goodfellow, Y Bengio, A Courville - 2016
- *Understanding machine learning: From theory to algorithms*, S Shalev-Shwartz, S Ben-David - 2014

Ce cours a été conçu avec Sandrine Mouysset et Axel Carlier.

## Representation of Perceptron



- 1 Inner product between input vector  $x \in \mathbb{R}^n$  and the weight  $w$  :  $w^T x$ ;
- 2 Add a bias scalar ( $b \in \mathbb{R}$ ) :  $z = w^T x + b$
- 3 Application of an activation function to  $z$  :  $a = f(z)$
- 4 Output value  $\hat{y} = a$ , e.g.  $\hat{y} \in \{0, 1\}$  for binary classification.

## Activation functions

The **activation functions**, denoted  $f$ , are usually non-linear functions. They can play a role of thresholding with 3 regimes,

- non-active: if the input value is under a threshold;
- transition phase: if the input value is close to the threshold;
- active: if the input value is above the threshold;





Figure: Sigmoid activation function :  $f(z) = \frac{1}{1+e^{-z}}$

Convert input vector to  $(x_1, \dots, x_n, 1)^T \in \mathbb{R}^{n+1}$

$$w^T x + b = (w_1, \dots, w_n, b)^T (x_1, \dots, x_n, 1)$$



- 1 Initial weight  $w^{(0)} = (w_i^{(0)})_{i \leq n}$
- 2 Draw training samples  $(x^{\{1\}}, y^{\{1\}}), \dots, (x^{\{m\}}, y^{\{m\}})$ .
- 3 Compute the output of Perceptron and a loss  $J(w)$ :

$$\hat{y}^{\{j\}}(w) = f\left(\sum_{i=1}^n w_i x_i^{\{j\}}\right) \text{ and } J(w) = \frac{1}{m} \sum_{j=1}^m \ell(\hat{y}^{\{j\}}(w), y^{\{j\}})$$

- 4 Update the weights from  $w^{(t)}$  to  $w^{(t+1)}$


$$w_i^{(t+1)} = w_i^{(t)} - \alpha^{(t)} \frac{\partial J}{\partial w_i}(w^{(t)})$$

where  $\alpha^{(t)}$  is a step size (learning rate)  $\alpha^{(t)} > 0$ .


- 5 Repeat 2-4 until convergence of  $w^{(t)}$  or  $J(w^{(t)})$ .

$\Rightarrow$  How to define the **cost function**  $\ell$  ?

## Classification and regression



Classification



Regression


**Classification (Logistic regression)** Assign a category to each observation

**Binary case :** false/true,  $y \in \{0, 1\}$ ,  $\hat{y} \in [0, 1]$


- sigmoid activation ( $\mathbb{R} \rightarrow [0, 1]$ ):  $f(z) = (1 + e^{-z})^{-1}$
- Loss function: logistic cost (cross-entropy):

$$\text{loss}(\hat{y}, y) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

## Classification and regression



Classification



Regression

**Linear Regression** Predict a real value of each observation :


- linear activation :  $f(z) = z$
- Mean squared error cost function (MSE):

$$\ell(\hat{y}, y) = (y - \hat{y})^2$$

⇒ How to solve this type of problem ?



**Gradient descent method: iterative method to find an optimal  $w^*$**




- ① Let  $\hat{y}(w) = f(\sum_{i=1}^n w_i x_i)$  and  $J(w) = \ell(\hat{y}(w), y)$
- Assume  $f$  is sigmoid.
  - What is the gradient of  $J(w)$  with respect to  $w_i$  in the following 2 cases?
    - $\ell(\hat{y}, y) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$
    - $\ell(\hat{y}, y) = (y - \hat{y})^2$
- ② Classification of a training set in 2d ( $n = 2$ ) with  $m = 4$  samples:

$x_1$	$x_2$	$y$
2	1	1
0	-2	1
-2	1	0
0	2	0


- Represent the data in 2d plane.
- Propose a weight  $w = (w_1, w_2)$  (without bias) such that the perceptron separates the training set with the sigmoid activation.
- What is the minimal  $J(w)$  with the logistic cost  $\ell$ ?

# Multi-layer perceptron and Multi-class classification



A multi-layer perceptron (MLP) is composed of an **input** layer, several **hidden** layers and an **output** layer.

The **depth** of the network above is  $L = 4$  (3 hidden layers plus one output layer).




Multi-layer perceptron :

- 1 Functionality
- 2 Interpretation
- 3 Activation function
- 4 Multi-class classification loss

In order to train a multi-layer perceptron, we need to understand the following computational steps:

- ➊ **Forward propagation** of input data to output;
- ➋ Compute a **loss** from the output;
- ➌ **Back propagation**: compute **gradients** of the loss with respect to the weights of the **output layer** and **hidden layers**;
- ➍ **Update** all the weights based on optimization methods.



The weights of layer  $k$ :  $W_{i,j}^{[k]}$  and  $b_i^{[k]}$ ,  $i$  output index,  $j$  input index. For depth  $L$ , we denote all the weights by  $\theta = (W^{[k]}, b^{[k]})_{k \leq L}$ , e.g.  $L = 2$

$$\hat{y}(x, \theta) = f \circ f^{[2]} \left( W^{[2]} f^{[1]} (W^{[1]} x + b^{[1]}) + b^{[2]} \right)$$

For an input  $x^{\{i\}}$ , we write the output  $\hat{y}^{\{i\}}(\theta) = \hat{y}(x^{\{i\}}, \theta)$

2) Compute the objective function after the forward-propagation:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \ell(y^{\{i\}}, \hat{y}^{\{i\}}(\theta))$$


3) Back-propagation: to compute the gradients  $\nabla_{\theta} J = (\frac{\partial J}{\partial \theta})^T$  from output to input by the *chain rule* in Calculus, e.g.

$$\nabla_{\theta} J = \frac{1}{m} \sum_{i=1}^m \left( \frac{\partial \hat{y}^{\{i\}}}{\partial \theta} \right)^T \nabla_{\hat{y}^{\{i\}}} \ell(y^{\{i\}}, \hat{y}^{\{i\}})$$

- Step 1: compute  $\nabla_{\hat{y}^{\{i\}}} \ell(y^{\{i\}}, \hat{y}^{\{i\}})$  for  $1 \leq i \leq m$ .
- Step 2: compute  $\frac{\partial \hat{y}^{\{i\}}}{\partial \theta}$  for  $1 \leq i \leq m$ .
- Step 3: compute  $\nabla_{\theta} J$ .

Step 1 can be solved as in the previous Quiz. [How about Step 2?](#)






Assume  $\hat{y} = f(a^{[2]}) \in \mathbb{R}$ ,  $a^{[2]} = f^{[2]}(W_{1,1}^{[2]}a_1^{[1]} + W_{1,2}^{[2]}a_2^{[1]} + b_1^{[2]}) \in \mathbb{R}$ .

Compute  $\frac{\partial \hat{y}}{\partial \theta}$ :

$$\frac{\partial \hat{y}}{\partial b_1^{[2]}} = \frac{\partial \hat{y}}{\partial a^{[2]}} \frac{\partial a^{[2]}}{\partial b_1^{[2]}} = f'(a^{[2]}) f^{[2]'}(W_{1,1}^{[2]}a_1^{[1]} + W_{1,2}^{[2]}a_2^{[1]} + b_1^{[2]})$$




Assume  $\hat{y} = f(a^{[2]}) \in \mathbb{R}$ ,  $a^{[2]} = f^{[2]}(W_{1,1}^{[2]}a_1^{[1]} + W_{1,2}^{[2]}a_2^{[1]} + b_1^{[2]}) \in \mathbb{R}$ .

Compute  $\frac{\partial \hat{y}}{\partial \theta}$ :

$$\frac{\partial \hat{y}}{\partial W_{1,1}^{[2]}} = \frac{\partial \hat{y}}{\partial a^{[2]}} \frac{\partial a^{[2]}}{\partial W_{1,1}^{[2]}} = f'(a^{[2]}) f^{[2]'}(W_{1,1}^{[2]}a_1^{[1]} + W_{1,2}^{[2]}a_2^{[1]} + b_1^{[2]}) a_1^{[1]}$$

# Illustration of back propagation




Assume  $\hat{y} = f(a^{[2]}) \in \mathbb{R}$ ,  $a^{[2]} = f^{[2]}(W_{1,1}^{[2]}a_1^{[1]} + W_{1,2}^{[2]}a_2^{[1]} + b_1^{[2]}) \in \mathbb{R}$ .

Compute  $\frac{\partial \hat{y}}{\partial \theta}$ : Assume  $a^{[1]} = f^{[1]}(W^{[1]}a^{[0]} + b^{[1]})$

$$\frac{\partial \hat{y}}{\partial W_{i,j}^{[1]}} = \frac{\partial \hat{y}}{\partial a^{[2]}} \frac{\partial a^{[2]}}{\partial a^{[1]}} \frac{\partial a^{[1]}}{\partial W_{i,j}^{[1]}}$$

Jacobian matrices:  $\frac{\partial a^{[1]}}{\partial W_{i,j}^{[1]}} : \mathbb{R}^1 \rightarrow \mathbb{R}^2$ ,  $\frac{\partial a^{[2]}}{\partial a^{[1]}} : \mathbb{R}^2 \rightarrow \mathbb{R}^1$ ,  $\frac{\partial \hat{y}}{\partial a^{[2]}} : \mathbb{R} \rightarrow \mathbb{R}$ ,

# Illustration of back propagation



Assume  $\hat{y} = f(a^{[2]}) \in \mathbb{R}$ ,  $a^{[2]} = f^{[2]}(W_{1,1}^{[2]}a_1^{[1]} + W_{1,2}^{[2]}a_2^{[1]} + b_1^{[2]}) \in \mathbb{R}$ .

Compute  $\frac{\partial \hat{y}}{\partial \theta}$ : Assume  $a^{[1]} = f^{[1]}(W^{[1]}a^{[0]} + b^{[1]})$

$$\frac{\partial a^{[2]}}{\partial a_1^{[1]}} = f^{[2]'}(W_{1,1}^{[2]}a_1^{[1]} + W_{1,2}^{[2]}a_2^{[1]} + b_1^{[2]})W_{1,1}^{[2]}$$

$$\frac{\partial a_i^{[1]}}{\partial W_{i,j}^{[1]}} = f_i^{[1]'}(W^{[1]}a^{[0]} + b^{[1]})a_j^{[0]}, \quad i, j \in \{1, 2\}$$






Figure: Linear vs. non-linear separation of training data

The non-linearity  $f^{[1]}, f^{[2]}, \dots$  in MLP plays a key role for non-linear separation.

Example: <https://playground.tensorflow.org/>




- The gradient function tends to zero when  $z$  is away from 0: cause vanishing gradients in the back propagation.
- Exercise: What is the relation between the sigmoid activation and the tanh activation:  $\text{Tanh}(z) = (e^z - e^{-z}) / (e^z + e^{-z})$ .




$$\text{ReLU}(z) = \begin{cases} z & \text{if } z > 0 \\ 0 & \text{if } z \leq 0 \end{cases}$$

The gradient is either 0 or 1 (when  $z \neq 0$ ), very different to that of the sigmoid.



$$\text{leakyReLU}(z) = \begin{cases} \alpha z & \text{if } z < 0 \\ z & \text{if not} \end{cases}$$

*Leaky Rectified Linear Unit*



$$\text{eLU}(z) = \begin{cases} \alpha(e^z - 1) & \text{if } z < 0 \\ z & \text{if not} \end{cases}$$

*Exponential Linear Unit*

These functions could potentially improve the gradient-descent training, e.g. to achieve faster convergence, but it is quite empirical.



- Question 1: Assume  $y \in \{1, 2, \dots, C\}$ . To classify  $x$  into  $C$  categories, how to design a differentiable loss  $\ell(y, \hat{y})$ ?
- Question 2: Assume  $\hat{y}$  is a probability distribution over  $\{1, 2, \dots, C\}$ , how to compute it from the output of an MLP?

- Answer 1: Use cross-entropy loss by representing  $y$  and  $\hat{y}$  as a probability distribution over  $\{1, 2, \dots, C\}$ .
- Answer 2: Design a non-linear differentiable function  $f$  such that  $\hat{y} = f(a)$ .

- KL divergence between two distributions  $p$  and  $q$  over  $\{1, \dots, C\}$

$$KL(q||p) = \sum_{i=1}^C \log \frac{q_i}{p_i} q_i$$

- Let  $y$  be a vector in  $\{0, 1\}^C$  such that  $y_i = 1$  i.f.f the category of  $y$  is  $i$ .
- Let  $\hat{y}$  be a vector in  $[0, 1]^C$  such that  $\sum_i \hat{y}_i = 1$ .
- The cross-entropy loss is  $KL(y||\hat{y})$ , which is equivalent to

$$\sum_{i=1}^C \log(y_i) y_i - \sum_{i=1}^C \log(\hat{y}_i) y_i$$

- In practice, we minimize the second term (to optimize MLP):

$$- \sum_{i=1}^C \log(\hat{y}_i) y_i$$

- When  $C = 2$ , this loss is equivalent to the logistic regression loss.

- Let  $a \in \mathbb{R}^C$  and  $\hat{y} = f(a)$ , such that

$$\hat{y}_i = \frac{e^{a_i}}{\sum_k e^{a_k}},$$

- We have  $\hat{y}_i \geq 0$  and  $\sum_i \hat{y}_i = 1$ .
- **Proposition:** for any  $c \in \mathbb{R}$ ,  $f(a + c) = f(a)$ .
- To avoid numerical issues, we compute  $f(a + c)$  with  $c = -\max_j a_j$ ,

$$\hat{y}_i = \frac{e^{a_i - \max_j a_j}}{\sum_k e^{a_k - \max_j a_j}}.$$

- Exercise: Compute the derivative  $\frac{\partial \hat{y}_i}{\partial a_k}$  for  $i = k$  and  $i \neq k$ .
- Exercise (TP): MNIST classification with MLP.